



Synergistic Algorithmic Repair Framework for Safeguarding Evidence-Based Information in Generative AI

Steven W. Giovinco*

AI Reputation Management, Recover Reputation/SW Giovinco Consulting LLC, 10003 New York, United States

* Correspondence: Steven W. Giovinco (steve@recoverreputation.com);

Received: 04-05-2025

Revised: 05-25-2025

Accepted: 06-01-2025

Citation: Giovinco, S. W. (2025). Synergistic Algorithmic Repair Framework for safeguarding evidence-based information in Generative AI. *J. Organ. Technol. Entrep.*, 3(2), 96–104. <https://doi.org/10.56578/jote030202>.



© 2025 by the author(s). Published by Acadlore Publishing Services Limited, Hong Kong. This article is available for free download and can be reused and cited, provided that the original published version is credited, under the CC BY 4.0 license.

Abstract: The proliferation of Large Language Models (LLMs) presents systemic challenges, including the generation of misinformation, hallucinations, and amplification of societal biases, causing a broader “epistemic crisis”. This paper introduced the novel Synergistic Algorithmic Repair Framework designed to mitigate the algorithmic harm. The methodology integrated three components into a continuous feedback loop: (1) The proactive curation of a verifiable digital ecosystem to establish ground truth; (2) the systematic application of Verifiable Human Feedback to correct LLM output with evidence-based inputs; and (3) the strategic curation of verified data into datasets for long-term model refinement. The efficacy of the framework was assessed through a qualitative case study involving two distinct real-world applications. The results demonstrated that the intervention successfully suppressed the targeted disinformation campaign of search engines and transformed LLM-generated output from negative or non-existent to positive or factual, in order to align with the curated ground truth. The study concluded that the synergistic operation of these components provided a durable and resilient method for repairing algorithmic harm. This framework, as a practical tool, offered promising insights for organizations and individuals to achieve digital equity and served as an operational model for implementing principles of ethical AI governance, such as human oversight and accountability.

Keywords: Generative AI; Algorithmic repair; Online reputation management; Digital ecosystem; Reinforcement learning from human feedback; Algorithmic injustice; Digital equity; AI governance

1. Introduction

The widespread societal integration of Large Language Models (LLMs) and other Generative AI systems has marked a paradigm shift in information retrieval, content creation, and decision support (Dwivedi et al., 2021). However, their proliferation presents profound and systemic challenges. Core failure modes inherent in current generation models incur the phenomenon of “hallucinations”, in which an LLM generates fluent and authoritative-sounding responses that are entirely fabricated or factually incorrect (Simon et al., 2023). Furthermore, because these models are trained on vast uncensored corpora of public web data, they invariably inherit and amplify biases, stereotypes, and discriminatory narratives present in that data, leading to “algorithmic injustice” (Benjamin, 2019; Noble, 2018; O’Neil, 2016). The ability of these systems to generate high-quality and persuasive misinformation at an unprecedented scale and speed poses a direct threat to information integrity (AbuJarour et al., 2024).

The consequences of these failures extend beyond isolated instances of incorrect information. The systematic distortion of information environments by algorithmically amplified falsehoods erodes public trust in institutions, fosters polarization, and degrades democratic decision-making (Lee-Geiller, 2024; Pew Research Center, 2024). This has led to a broader societal issue characterized as an “epistemic crisis”, where a general loss of trust in all forms of information becomes pervasive, hence increasing the difficulty to establish a shared factual basis for public discourse (Lăzăroiu et al., 2020). This crisis represents a significant technical problem requiring a new class of solutions capable of re-establishing and enforcing verifiable truth within a corrupted information ecosystem. It has been proved that traditional countermeasures are inadequate to address the scale and nature of this problem.

Post-hoc interventions such as debunking, source inoculation, and labeling of AI-generated content have demonstrated limited effectiveness. Research indicated that simple disclaimers or warning labels often failed to reduce user reliance on misinformation and might induce a general and indiscriminating skepticism rather than fostering critical judgment (Altay et al., 2023).

Similarly, conventional approaches in Online Reputation Management (ORM) are fundamentally mismatched to address the challenge of LLM-generated harm. Traditional ORM is a discipline focused primarily on influencing the ranking of content on Search Engine Results Pages (SERPs) through techniques such as Search Engine Optimization (SEO), link suppression, and content generation (McCombs & Shaw, 1972). These strategies are inherently reactive, to address damage to reputation after it has manifested on the public web (Coombs, 2014). Crucially, they lack any mechanism for direct intervention or repair of the underlying AI models that are gradually becoming the source of this damage. This fundamental difference in approaches is illustrated in Figure 1.




Figure 1. Comparison of the traditional siloed approach to algorithmic repair with the proposed integrated and synergistic framework

There exists a fundamental paradigm mismatch between the operation of traditional ORM and the architecture of LLMs. Conventional ORM operates on the “presentation layer” of the internet, i.e., the SERP, aiming to control which pre-existing documents are most visible. In contrast, the proposed framework operates on the “knowledge layer”; in other words, the internal parametric state of LLM and its underlying training data. The output of LLM is not the exclusive result of a real-time web search but is generated from its pre-trained knowledge base. The primary goal of traditional ORM is to alter the top ten search results for a given query; this has a slow, indirect, and often negligible effect on the output of an already-trained LLM. The model is not re-indexing the web for each query, rendering the tools of conventional ORM ineffective at addressing the root cause of AI-generated misinformation.

Consequently, there exists a critical and unaddressed gap in the literature for a proactive, integrated, and synergistic framework that must bridge the divide between the strategic creation of authoritative content on the web and the evidence-based use of that content to repair and refine the generative models themselves. This paper addressed this gap by proposing and evaluating the novel Synergistic Algorithmic Repair Framework. Using qualitative case studies as a methodology, this research demonstrated how the integrated approach of the framework provided a durable and effective means to mitigate AI-generated misinformation and promote digital equity. The following sections will detail the methodology of the framework, present the results of its application in two case studies, and discuss the broader implications for AI governance and practice.

2. Methodology

Qualitative case study, a research approach chosen for its suitability, could help explore a contemporary phenomenon in the real-life context. To enhance methodological rigor and mitigate potential practitioner-researcher bias, a process of data triangulation was employed to validate outcomes. This research strategy involved using multiple data sources to develop a comprehensive and corroborated understanding of a phenomenon. Validating the effectiveness of the framework required systematic collection and comparison of data from: (1) Direct LLM Output; (2) Search Engine Results Page (SERP) Analysis; and (3) Third-Party Media and Platform Analysis. Reliance on a single and potentially biased observation was reduced by cross-verifying findings across these distinct data sources, thereby strengthening the validity of the results.

The Synergistic Algorithmic Repair Framework, which operates as a continuous feedback loop (Figure 2), serves as an intervention protocol being investigated. The framework begins with a proactive auditing phase, which serves as a systematic diagnosis of the digital footprint of an entity. This audit is conducted across both traditional search engines and a plurality of LLM platforms, e.g., OpenAI, ChatGPT, Google Gemini, and Anthropic Claude,

to identify informational inaccuracies, including factual errors, propagation of negative or outdated narratives, harmful biases, and “information vacuums” where the model lacks sufficient data to provide a response.




Figure 2. Synergistic and continuous feedback loop of the algorithmic repair framework, illustrating the interplay of the three core pillars

2.1 Data Collection and Analysis

This audit is a quantitative and data-driven process that employs reproducible metrics to assess the representation of an entity within AI-generated output. These metrics include:

- **Citation Sentiment Score:** A weighted score derived by systematically collecting branded citations from LLM output and tagging each mention as positive, neutral, or negative. This provides a quantifiable measure for the overall tone of the AI portrayal.
- **Source Trust Differential:** An analysis of the sources cited by LLM in its responses. Each source is assigned a trust score based on its authority (e.g., Tier 1 for high-authority global media, Tier 2 for niche publications, and Tier 3 for forums and unverified blogs). This metric evaluates the quality of the information ecosystem the LLM is referencing.
- **Narrative Consistency Index:** A measure of the alignment between the key messages in LLM-generated responses and the core and strategic narrative of an entity. This identifies instances where AI is misrepresenting or diluting the intended positioning of the entity.
- **Entity Co-Occurrence Mapping:** An analysis that identifies other entities (e.g., competitors, partners, and concepts) that are frequently mentioned alongside the subject entity in LLM output. These co-occurrences are categorized as positive, neutral, or negative associations to map the contextual reputations being constructed by AI. The output of this diagnostic audit is a detailed map of the specific informational inaccuracies that must be corrected by subsequent phases of the framework.

2.2 Pillar 1: Digital Ecosystem Curation (Ground Truth Generation)

This foundational pillar is responsible for the strategic creation of a corpus of verifiable and high-quality source materials that will serve as the “ground truth” for all subsequent repair activities. This process is explicitly distinguished from conventional ORM. While it employs ORM techniques for discoverability, its primary technical objective is to generate a corpus of data that is optimized for ingestion, comprehension, and validation by AI systems, not merely to achieve high rankings on a search results page. The steps in this phase include:

1. **Develop a Comprehensive Blueprint:** Based on the data-driven audit, a strategic content and platform plan are developed to directly and verifiably counter each identified inaccuracy.
2. **Build and Strengthen Authoritative Platforms:** This involves creating or optimizing high-authority digital assets, such as dedicated corporate or personal websites, professional profiles (e.g., LinkedIn), and entries on community-driven knowledge platforms (e.g., Wikipedia and Reddit). A critical technical step is the implementation of structured data using schema markup, which provides explicit contextual information to help search engines and AI systems better understand and categorize the content, thereby improving the accuracy of their knowledge extraction processes.
3. **Publish High-Quality and Relevant Content:** A library of factually accurate and thoroughly researched content is created and published on high-authority third-party platforms, which are significant components of LLM training corpora. This content may include white papers, peer-reviewed articles, thought-leadership pieces, and multimedia assets designed to demonstrate expertise. It provides credible and citable sources that directly address and correct identified misinformation.

2.3 Pillar 2: Verifiable Human Feedback Integration (Direct Intervention)

This pillar leverages the built-in feedback mechanisms of LLM platforms as a tool for direct and targeted intervention. The process is aligned with the principles of Reinforcement Learning from Human Feedback (RLHF), a technique with which human evaluations are used to guide a model’s behavior toward more accurate, ethical, and helpful output (Christiano et al., 2017). The key methodological step is the introduction of verifiability into this feedback process. The human feedback provided is neither subjective nor based on personal opinions. Instead, evaluators are instructed to systematically provide corrective inputs that are explicitly and demonstrably based on the verifiable source materials created in Pillar 1. When LLM generates an information inaccuracy, corrective feedback provided through the user interface comprises not only the corrected factual information but also a direct citation or URL linking to the authoritative source document within the curated digital ecosystem. This transforms a generic feedback feature into a precise and evidence-based algorithmic repair mechanism, to create a traceable link between the corrected output of the model and the ground truth.

2.4 Pillar 3: Strategic Dataset Curation (Long-Term Inoculation)

The third pillar ensures the long-term durability and resilience of algorithmic repair. The corpus of verifiable and authoritative content, generated in Pillar 1 and validated in the feedback process of Pillar 2, is systematically collected, structured, and formatted into a high-quality dataset. This curated dataset serves as a long-term solution for model refinement. It can be simply provided to the developers of LLM for inclusion in subsequent training runs, to ensure that future versions of the model are trained on a factually accurate and unbiased representation of the entity from the outset. Alternatively, the dataset can be used to fine-tune a proprietary or specialized instance of LLM. This process directly addresses the documented need in the field for diverse, high-quality, and curated datasets to mitigate systemic bias and improve the baseline accuracy of AI models over time.

3. Results

Case studies A and B, with application of the framework, consistently followed the model of transformation as illustrated in Figure 3. A comprehensive comparative analysis detailing the specific contexts, interventions, and outcomes for each case study is provided in Table 1.




Figure 3. Model of the intervention process, showing the transformation of an AI’s output from an initial detrimental state to a beneficial state post-intervention

Table 1. Comparative analysis of case studies A and B, after having implemented the framework

Feature	Case Study A (Hedge Fund CEO)	Case Study B (Sustainable Energy Group)
Context	Reactive: Individual defamation campaign	Proactive: Organizational misinformation
Web issue	Five defamatory articles on first-page search results.	Six negative articles on first-page search results across multiple countries.
AI issue	“Information vacuum” on LLM platform (Google Gemini).	Propagation of false narratives by LLM (ChatGPT).
Primary intervention	ORM to address information vacuum and suppress negative content.	Multilingual ORM to establish a positive and factual narrative.
AI strategy	Direct feedback and targeted content seeding.	Systematic feedback and corrective dataset provision.
Timeline	6 months	6 months
ORM results	100% suppression of negative search results.	100% suppression of negative search results.
AI results	LLM output transformed from “information vacuum” to positive summary.	LLM output corrected from negative narrative to positive summary.
Overall impact	Neutralized defamation campaign and protected financial reputation.	Countered misinformation and converted AI into a brand-building asset.

Note: LLM = Large Language Model; ORM = Online Reputation Management.

3.1 Case Study A: Correcting an Information Vacuum

The framework was applied to the case of a hedge fund CEO facing a targeted smear campaign. The initial audit revealed five defamatory articles on the first page of Google search results and an “information vacuum” on Google Gemini, which returned no information about the executive (Figure 4).




Figure 4. Simplified representation of an “information vacuum”, indicating the state of AI output on LLM prior to intervention in case study A
Note: LLM = Large Language Model.

The intervention involved a six-month application of the framework where a dedicated personal website was built and profiles on niche financial platforms were optimized. High-quality financial articles were published on authoritative platforms to serve as credible sources. A direct feedback loop with Gemini was initiated to reinforce the accuracy of the new content. Within four months, complete suppression of the defamatory content from the first page of search results was observed. Concurrently, the output of Gemini was transformed from providing no information to generating a positive and factually accurate summary of the CEO’s career, drawing directly from the newly created content (Figure 5).




Figure 5. Simplified representation of the transformed LLM output for case study A after the intervention, with which a positive and factually accurate summary was projected
Note: LLM = Large Language Model.

3.2 Case Study B: Countering a Disinformation Campaign

The framework was applied to a sustainable energy company facing a disinformation campaign, resulting in six negative links on Google and propagation of these falsehoods by ChatGPT. A multilingual strategy was deployed. The Digital Ecosystem Curation phase involved enhancing the company’s corporate website and Wikipedia entries with verified information via publishing a library of accurate content on high-authority platforms.

Concurrently, the feedback system of ChatGPT was methodically used to report false narratives while providing links to new and authoritative content. Within six months, the negative search content was suppressed in all

targeted countries. In parallel, the output of ChatGPT shifted from providing damaging information to generating a detailed and positive summary of the leadership (Figure 6).




Figure 6. Simplified representation of the correct LLM output for case study B on a different platform, which successfully established a positive narrative
Note: LLM = Large Language Model.

4. Discussion

Results from the case studies indicated that the Synergistic Algorithmic Repair Framework was effective in repairing targeted algorithmic harm. The success of the interventions supported the central hypothesis that a synergistic approach, integrating Digital Ecosystem Curation, verifiable human feedback, and AI ethics (Figure 7), was required to produce durable and meaningful changes in LLM output. This integrated nature stood in contrast to the siloed and reactive approach of conventional methods, as evidenced by the comparative analysis in Table 2.




Figure 7. Conceptual foundation of the synergistic framework, illustrating its integration of principles from digital ecosystem curation, Reinforcement Learning from Human Feedback, and AI ethics

Table 2. Comparison of algorithmic repair methodologies

Feature	Traditional SEO-Based ORM	Disconnected AI Feedback	The Synergistic Framework
Scope of action	Web search results (presentation layer)	Single LLM output (instance layer)	Entire digital ecosystem (knowledge layer)
Mechanism	Reactive content suppression/promotion	Subjective and ad-hoc user corrections	Proactive, evidence-based, and systematic repair
Ground truth	Not explicitly defined	User's personal knowledge or opinion	Explicitly created a corpus of verifiable content
Durability of the effect	Temporary; vulnerable to algorithm changes	Limited; may not persist or generalize	Durable and self-reinforcing
Ethical alignment	Prone to misuse (e.g., astroturfing)	Neutral; lacked a guiding ethical framework	Designed for digital equity and factual accuracy

Note: SEO = Search Engine Optimization; ORM = Online Reputation Management; LLM = Large Language Model.

4.1 Implications for Practice and Theory

The findings had broader implications for both practice and theory. The framework could be framed as a practical tool for promoting digital equity and inclusion. It provided a structured method for individuals and marginalized communities to challenge and repair the “algorithmic injustice” that could arise from biased training data, to ensure fairer and more accurate representation in AI-driven information ecosystems (Kinchin, 2024).

Furthermore, the framework served as a tangible and operational embodiment of high-level AI ethical principles (Singla et al., 2025). It put these principles into practice:

- **Human Oversight:** The framework operationalizes human oversight by ensuring that human-verified and evidence-based information is the ultimate arbiter in the AI refinement process (UNESCO, 2021).
- **Fairness and Non-Discrimination:** It provides a direct mechanism to identify and correct harmful stereotypes and biases.
- **Transparency and Accountability:** It enhances transparency by creating a traceable and auditable link between a corrected AI output and the public, as well as verifiable source materials to justify the correction.

In addition to operationalizing these core ethical principles, the proposed framework demonstrates practical alignment with major international AI governance models. For instance, the EU AI Act categorizes AI systems by risk and mandates transparency and human oversight for high-risk applications—requirements directly supported by this framework’s emphasis on verifiable data and human-led feedback (European Parliament & Council of the European Union, 2024; Ismail & Ahmad, 2025). Likewise, the OECD AI Principles call for human-centered values, fairness, transparency, and accountability; all of which are operationalized through the synergistic pillars of Digital Ecosystem Curation, Verifiable Human Feedback, and Strategic Dataset Curation (OECD, 2024). By translating these high-level principles into a concrete, repeatable methodology, the framework offers organizations a tangible pathway toward compliance with emerging global standards for responsible AI.

4.2 Ethical Risks and Safeguards

The application of any reputation management framework necessitates a deep engagement with its ethical implications. A significant risk is the potential for misuse, such as “reputation laundering”, in which powerful individuals or organizations could employ such techniques to obscure unethical behavior or silence legitimate criticism. The Synergistic Algorithmic Repair Framework was designed with ethical safeguards to mitigate these risks. Unlike unethical ORM practices that may rely on creating fake reviews/astroturfing, censoring complaints, or disseminating false information (Campbell, 2025), this framework was fundamentally grounded in the principle of establishing a verifiable and evidence-based “wall of truth”. The core methodology required that all interventions from content creation to direct AI feedback were based on factually accurate, high-quality, and citable source materials. This emphasis on transparency and verifiability served as a crucial ethical guardrail, aligning the framework with principles of honest and authentic communication. The objective is not to erase history or silence dissent, but to ensure that the digital narrative is built upon a foundation of demonstrable truth (Entman, 1993).

5. Conclusions

This research introduced and provided empirical support for the Synergistic Algorithmic Repair Framework, an integrated system that could effectively mitigate AI-generated harm and establish a durable and factually accurate digital narrative. It is important to acknowledge the limitations of the present study and to outline directions for future research. The primary limitation is its reliance on two case studies for providing valuable proof-of-concept, thus restricting the generalizability of the findings. Owing to the contextual limitations of the current study, the outcomes observed may not be universally applicable across all industries, geographies, or types of reputational challenges. Future studies should aim at cross-industry validation to test the applicability of the framework in diverse contexts, such as non-profit, governmental, and small business sectors. To extend on this work, several avenues for further research are proposed:

- **Broader Empirical Testing:** Future research should focus on broader empirical testing of the framework across a diverse range of industries, varied by the sizes of organizations (e.g., non-profit making and small businesses) and geographical contexts. This would help validate its robustness and identify context-specific modifications that may be required.
- **Longitudinal Studies:** The long-term sustainability of the algorithmic repairs is a critical area for investigation. Longitudinal studies were designed to track the durability of the corrected AI output over extended periods; they could assess whether the positive narrative persists or continuous intervention is necessary to prevent regression.
- **Policy and Regulatory Implications:** Further research should explore the implications of this framework for policymakers and regulators. The model provides a tangible example of how platform accountability can be

operationalized. This could inform the development of regulations that encourage or mandate transparent and evidence-based feedback mechanisms for AI platforms, thus contributing to broader discussions on AI governance by multi-stakeholders and the global effort to combat disinformation.

Informed Consent Statement

Informed consent was obtained from all subjects involved in the study.

Data Availability

The data presented in this study are available on request from the corresponding author.

Acknowledgements

The author would like to acknowledge the contributions of the research team for their support.

Conflicts of Interest

The author declares no conflicts of interest.

References

- AbuJarour, S., Qarariah, A., Saadeh, N., & Salem, M. (2024). AI, misinformation, and fake news: A literature review of ethical and technical approaches. In *Finance and Law in the Metaverse World: Regulation and Financial Innovation in the Virtual World* (pp. 641–652). Cham: Springer. https://doi.org/10.1007/978-3-031-67547-8_55.
- Altay, S., Berriche, M., & Acerbi, A. (2023). A survey of expert views on misinformation: Definitions, determinants, solutions, and the future of the field. *Harv. Kennedy Sch. Misinformation Rev.*, 4(4). 1–34. <https://doi.org/10.37016/mr-2020-119>.
- Benjamin, R. (2019). *Race After Technology: Abolitionist Tools for the New Jim Code*. Polity Press.
- Campbell, K. (2025). *The ethics of online reputation management today*. <https://blog.reputationx.com/ethical-reputation-management>.
- Christiano, P. F., Leike, J., Brown, T., Martic, M., Legg, S., & Amodei, D. (2017). Deep reinforcement learning from human preferences. In *Advances in Neural Information Processing Systems 30, Long Beach Convention Center, Long Beach, CA, USA*. (pp. 4299–4307). <https://doi.org/10.48550/arXiv.1706.03741>.
- Coombs, W. T. (2014). *Ongoing Crisis Communication: Planning, Managing, and Responding (4th ed.)*. SAGE Publications.
- Dwivedi, Y. K., Hughes, L., Ismagilova, E., Aarts, G., Coombs, C., Crick, T., Duan, Y., Dwivedi, R., Edwards, J., Eirug, A., et al. (2021). Artificial intelligence (AI): Multidisciplinary perspectives on emerging challenges, opportunities, and agenda for research, practice and policy. *Int. J. Inf. Manage.*, 57, 101994. <https://doi.org/10.1016/j.ijinfomgt.2019.08.002>.
- Entman, R. M. (1993). Framing: Toward clarification of a fractured paradigm. *J. Commun.*, 43(4), 51–58. <https://doi.org/10.1111/j.1460-2466.1993.tb01304.x>.
- European Parliament & Council of the European Union. (2024). Regulation (EU) 2024/1689 of the European Parliament and of the Council of 13 June 2024 laying down harmonised rules on Artificial Intelligence (Artificial Intelligence Act). *Off. J. Eur. Union, L 2024/1689*. <https://eur-lex.europa.eu/eli/reg/2024/1689/oj/eng>.
- Ismail, O. & Ahmad, N. (2025). Ethical and governance frameworks for artificial intelligence: A systematic literature review. *Int. J. Interact. Mob. Technol.*, 19(14), 121–136. <https://doi.org/10.3991/ijim.v19i14.56981>.
- Kinchin, N. (2024). “Voiceless”: The procedural gap in algorithmic justice. *Int. J. Law Inf. Technol.*, 32, eaac024. <https://doi.org/10.1093/ijlit/eaac024>.
- Lăzăroiu, G., Neguriță, O., Grecu, I., Grecu, G., & Mitran, P. C. (2020). Consumers’ decision-making process on social commerce platforms: Online trust, perceived risk, and purchase intentions. *Front. Psychol.*, 11, 890. <https://doi.org/10.3389/fpsyg.2020.00890>.
- Lee-Geïller, S. (2024). The moderating effect of digital literacy on the link between e-government effectiveness and trust in government. *J. Policy Stud.*, 39(4), 77–101. <https://doi.org/10.52372/jps.e672>.
- McCombs, M. & Shaw, D. (1972). The agenda-setting function of mass media. *Public Opin. Q.*, 36(2), 176–187. <https://doi.org/10.1086/267990>.
- Noble, S. U. (2018). *Algorithms of Oppression: How Search Engines Reinforce Racism*. New York, NYU Press.
- OECD. (2024). *OECD AI Principles Overview*. <https://oecd.ai/en/ai-principles>.

- O'Neil, C. (2016). *Weapons of Math Destruction: How Big Data Increases Inequality and Threatens Democracy*. Crown Publishers.
- Pew Research Center. (2024). *Public Trust in Government: 1958–2024*. <https://www.pewresearch.org/politics/2024/06/24/public-trust-in-government-1958-2024/>
- Simon, F. M., Altay, S., & Mercier, H. (2023). Misinformation reloaded? Fears about the impact of generative AI on misinformation are overblown. *Harv. Kennedy Sch. Misinformation Rev.*, 4(5). 1–11. <https://doi.org/10.37016/mr-2020-127>.
- Singla, A., Sukharevsky, A., Yee, L., Chui, M., & Hall, B. (2025). *The state of AI: How organizations are rewiring to capture value*. McKinsey & Company. <https://www.mckinsey.com/capabilities/quantumblack/our-insights/the-state-of-ai>.
- UNESCO. (2021). *Recommendation on the ethics of artificial intelligence*. <https://unesdoc.unesco.org/ark:/48223/pf0000379920>.